

References

- [1] Swanson, D. R. (1986). *Undiscovered Public Knowledge*. **The Library Quarterly: Information, Community, Policy**, 56(2), 103–118. <https://doi.org/10.1086/601720>.
- [2] Langley, P., Simon, H. A., Bradshaw, G. L., & Zytkow, J. M. (1987). *Scientific Discovery: Computational Explorations of the Creative Processes*. MIT Press, Cambridge, MA. ISBN 978-0-262-12163-3.
- [3] Zhou, Y., Liu, H., Srivastava, T., Mei, H., & Tan, C. (2024). *Hypothesis Generation with Large Language Models*. In **Proceedings of the 1st Workshop on NLP for Science (NLP4Science)**, 117–139. Association for Computational Linguistics. DOI: 10.18653/v1/2024.nlp4science-1.10.
- [4] Yang, Z., Du, X., Li, J., Zheng, J., Poria, S., & Cambria, E. (2024). *Large Language Models for Automated Open-domain Scientific Hypotheses Discovery*. In **Findings of the Association for Computational Linguistics: ACL 2024**, 13545–13565. Association for Computational Linguistics. DOI: 10.18653/v1/2024.findings-acl.804.
- [5] Si, C., Yang, D., & Hashimoto, T. (2025). *Can LLMs Generate Novel Research Ideas? A Large-Scale Human Study with 100+ NLP Researchers*. In **International Conference on Learning Representations (ICLR 2025)**. arXiv:2409.04109. DOI: 10.48550/arXiv.2409.04109.
- [6] Du, Y., Li, S., Torralba, A., Tenenbaum, J. B., & Mordatch, I. (2024). *Improving Factuality and Reasoning in Language Models through Multiagent Debate*. In **Proceedings of the 41st International Conference on Machine Learning**, PMLR 235, 11733–11763. arXiv:2305.14325. DOI: 10.48550/arXiv.2305.14325.
- [7] Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhumoye, S., Yang, Y., Gupta, S., Majumder, B. P., Hermann, K., Welleck, S., Yazdanbakhsh, A., & Clark, P. (2023). *Self-Refine: Iterative Refinement with Self-Feedback*. In **Advances in Neural Information Processing Systems 36 (NeurIPS 2023)**. arXiv:2303.17651.
- [8] Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., & Yao, S. (2023). *Reflexion: Language Agents with Verbal Reinforcement Learning*. In **Advances in Neural Information Processing Systems 36 (NeurIPS 2023)**. arXiv:2303.11366. DOI: 10.48550/arXiv.2303.11366.
- [9] Hong, S., Zhuge, M., Chen, J., Zheng, X., Cheng, Y., Ceyao, Z., Wang, J., Wang, Z., Yau, S. K. S., Lin, Z., Zhou, L., Ran, C., Xiao, L., Wu, C., & Schmidhuber, J. (2024). *MetaGPT: Meta Programming for A Multi-Agent Collaborative Framework*. In **International Conference on Learning Representations (ICLR 2024)**. arXiv:2308.00352.
- [10] Smit, A. P., Grinsztajn, N., Duckworth, P., Barrett, T. D., & Pretorius, A. (2024). *Should we be going MAD? A Look at Multi-Agent Debate Strategies for LLMs*. In **Proceedings of the 41st International Conference on Machine Learning**, PMLR 235, 45883–45905. arXiv:2311.17371.
- [11] Skarlinski, M. D., Cox, S., Laurent, J. M., Braza, J. D., Hinks, M., Hammerling, M. J., Ponnampati, M., Rodrigues, S. G., & White, A. D. (2024). *Language Agents Achieve Superhuman Synthesis of Scientific Knowledge*. arXiv:2409.13740.
- [12] Gottweis, J., Weng, W.-H., Daryin, A., Tu, T., Palepu, A., Sirkovic, P., Myaskovsky, A., Weissenberger, F., Rong, K., Tanno, R., Saab, K., Popovici, D., Blum, J., Zhang, F., Chou, K.,

Hassidim, A., Gokturk, B., Vahdat, A., Kohli, P., Matias, Y., Carroll, A., Kulkarni, K., Tomasev, N., Guan, Y., Dhillon, V., Dhaval Vaishnav, E., Lee, B., Costa, T. R. D., Penadés, J. R., Peltz, G., Xu, Y., Pawlosky, A., Karthikesalingam, A., & Natarajan, V. (2025). *Towards an AI co-scientist*. arXiv:2502.18864. DOI: 10.48550/arXiv.2502.18864.

[13] Pitre, P., Ramakrishnan, N., & Wang, X. (2025). *CONSENSAGENT: Towards Efficient and Effective Consensus in Multi-Agent LLM Interactions Through Sycophancy Mitigation*. In **Findings of the Association for Computational Linguistics: ACL 2025**, 22112–22133. Association for Computational Linguistics. DOI: 10.18653/v1/2025.findings-acl.1141.

[14] Boiko, D. A., MacKnight, R., Kline, B., & Gomes, G. (2023). *Autonomous chemical research with large language models*. **Nature**, **624**, 570–578. DOI: 10.1038/s41586-023-06792-0.

[15] *The Virtual Lab: AI Agents Design New SARS-CoV-2 Nanobodies with Experimental Validation* (2024). bioRxiv preprint 2024.11.11.623004. DOI: 10.1101/2024.11.11.623004.

[16] Lu, C., Lu, C., Lange, R. T., Foerster, J., Clune, J., & Ha, D. (2024). *The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery*. arXiv:2408.06292. DOI: 10.48550/arXiv.2408.06292.

[17] Yamada, Y., Lange, R. T., Lu, C., Hu, S., Foerster, J., Ha, D., & Clune, J. (2026). *Towards End-to-End Automation of AI Research*. **Nature**. Published online 25 March 2026. DOI: 10.1038/s41586-026-10265-5. arXiv:2606.15497.

[18] Yamada, Y., Lange, R. T., Lu, C., Hu, S., Lu, C., Foerster, J., Clune, J., & Ha, D. (2025). *The AI Scientist-v2: Workshop-Level Automated Scientific Discovery via Agentic Tree Search*. arXiv:2504.08066.

[19] Goo, T., Kim, J., Han, K., Jo, G., Kim, J.-S., & Kim, T.-H. (2026). *Evaluating and Guarding Citation Faithfulness in Agentic Scientific Synthesis*. arXiv:2607.20527.

[20] Zhang, G., Yue, Y., Li, Z., Yun, S., Wan, G., Wang, K., Cheng, D., Yu, J. X., & Chen, T. (2025). *Cut the Crap: An Economical Communication Pipeline for LLM-Based Multi-Agent Systems*. In **International Conference on Learning Representations (ICLR 2025)**.

[21] *Scientific hypothesis generation by large language models: laboratory validation in breast cancer treatment* (2025). **Journal of the Royal Society Interface**. DOI: 10.1098/rsif.2024.0674. PMID: 40462712; PMCID: PMC12134935.

[22] Zhu, K., Zhang, J., Qi, Z., Shang, N., Liu, Z., Han, P., Su, Y., Yu, H., & You, J. (2025). *SafeScientist: Toward Risk-Aware Scientific Discoveries by LLM Agents*. arXiv:2505.23559.